

Combining Apache Kafka and the Elastic Stack

Apache Kafka

Brokers

Topics

Partitions

Consuming topics

Elastic Stack

Ingesting data into Elasticsearch

Visualizing data

Combining Apache Kafka and the Elastic Stack

Demo

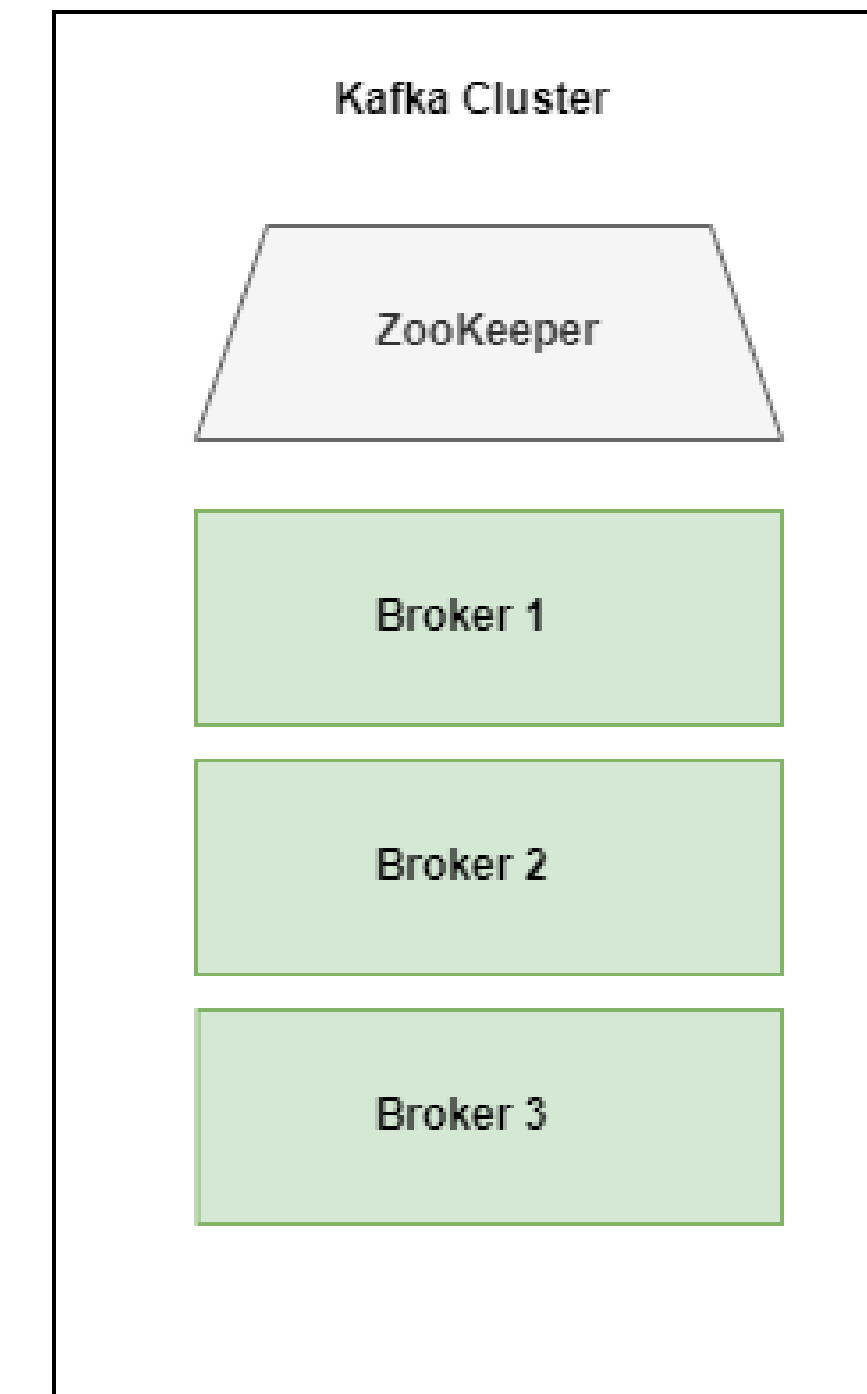
Apache Kafka

- Open-source distributed event streaming platform
- Provides high-throughput, low-latency
- Kafka Connect can be used to connect with external systems
 - Source Connector and Sink Connector
- REST APIs are also available for producing, consuming and streaming
- Stores data in event logs called **topics**
- Kafka is **pull-based**
 - Consumers are able to ask for new messages when they are ready
 - Makes messages replayable
- Available as both a fully-managed and self-managed service
 - Confluent Cloud is used in the demo



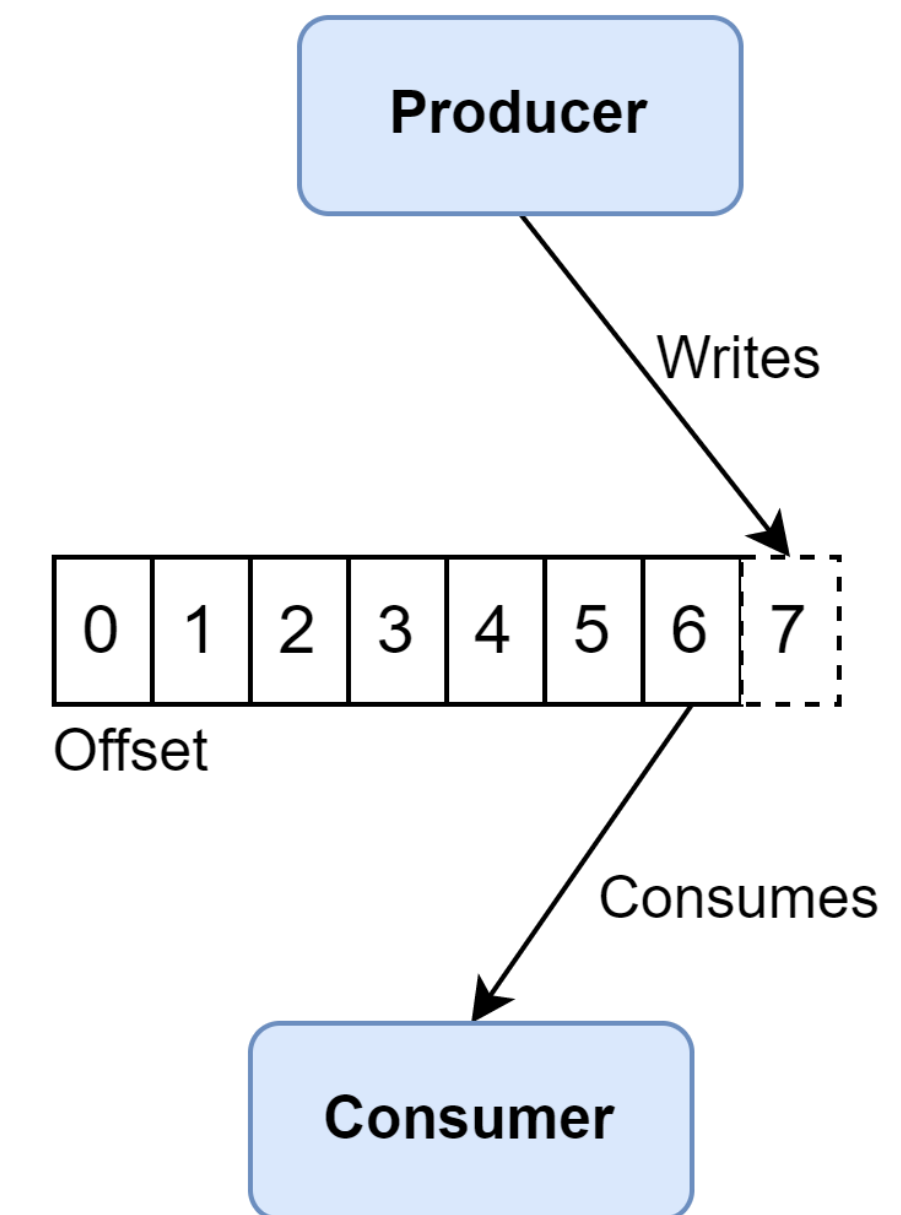
Brokers

- Kafka Brokers are servers that are part of a Kafka Cluster
- **Topics** are stored in brokers
- A Kafka cluster typically consists of several brokers
 - This gives the benefit of data replication as topics are replicated across multiple brokers
- Kafka brokers are stateless, and uses Apache ZooKeeper for maintaining their cluster state



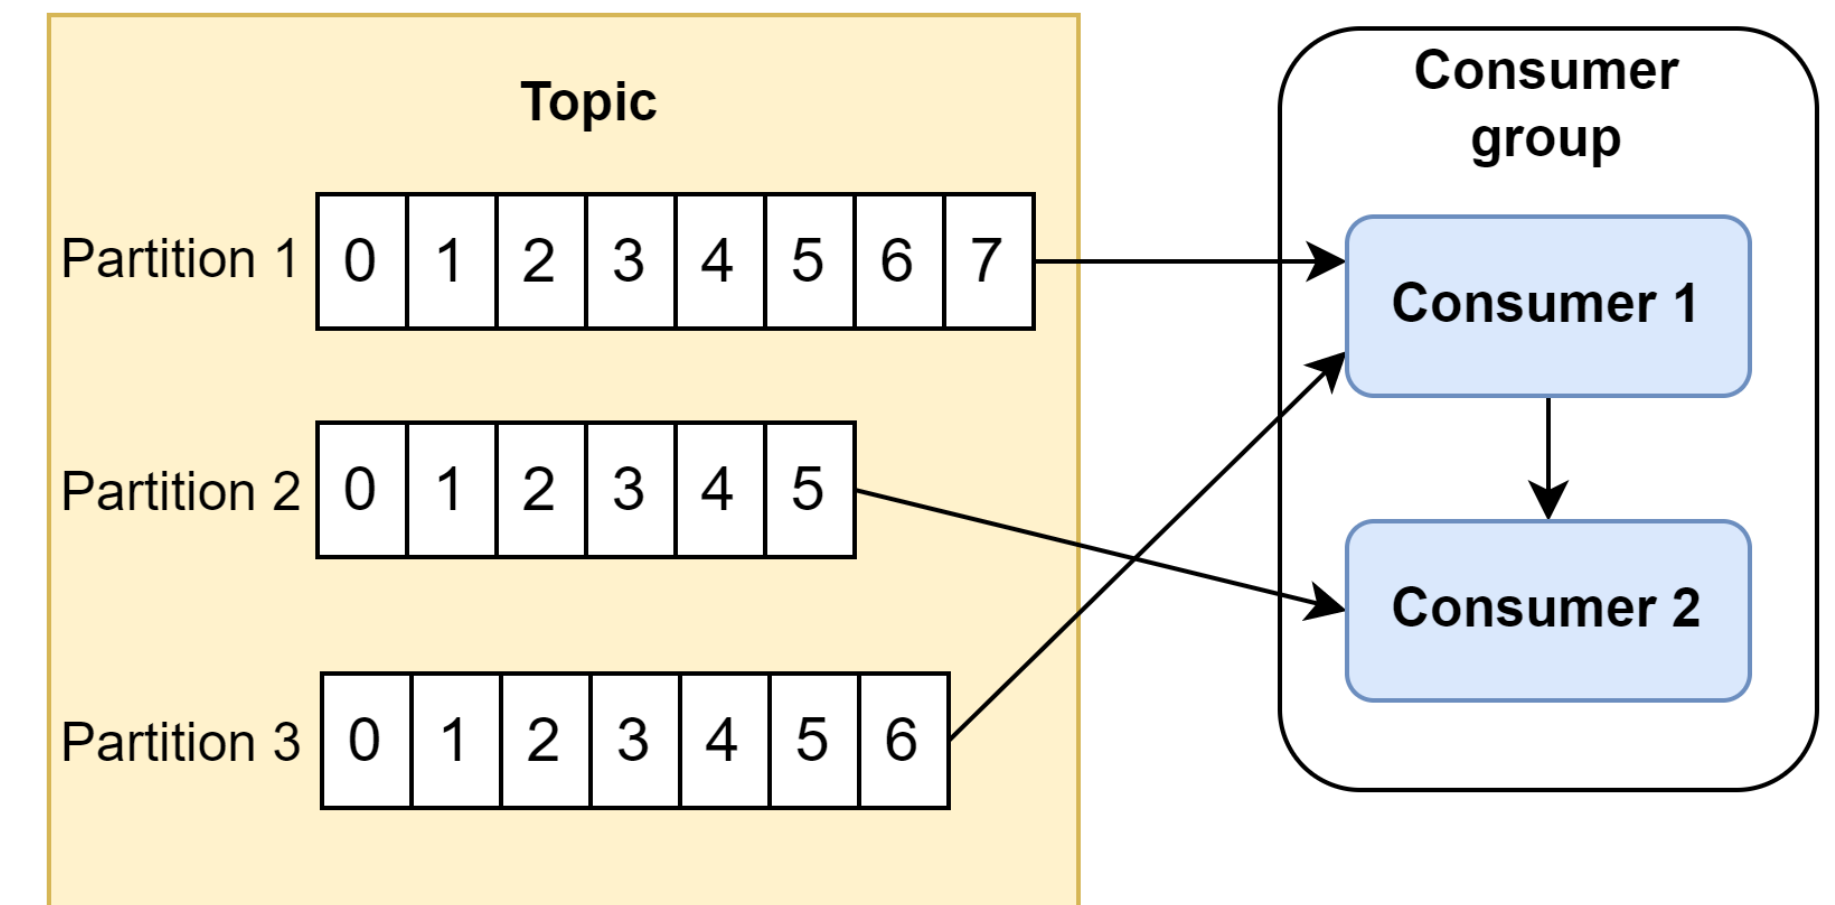
Topics

- Log containing an ordered collection of messages
- Each topic is append-only, and a message is given an incremental **offset** upon storage
 - A message consists of a key, value, offset, timestamp and headers
- Topics can have multiple **producers** and **consumers**
- A topic can store data anywhere from a short amount of time, to indefinitely
 - This is decided upon creation by the chosen retention policy
 - Upon expiry, messages are marked for deletion
 - Compaction policy retains only the most recent message for each key
- **Schema Registry** allows us to enforce schemas on messages
- Topics are divided into **partitions**



Partitions

- The smallest storage unit in Kafka
- Each partition holds a subset of messages in a topic
- A topic can have one or several partitions
 - More partitions allows for more parallelism
- The offset of a message is based on its partition
 - The offset guarantees the order within a partition, but not across the topic
- A message can be assigned to a specific partition by specifying a **partition key**
 - If no partition key is provided, Kafka will use round-robin assignment
- Each partition is assigned to exactly one consumer within a **consumer group**

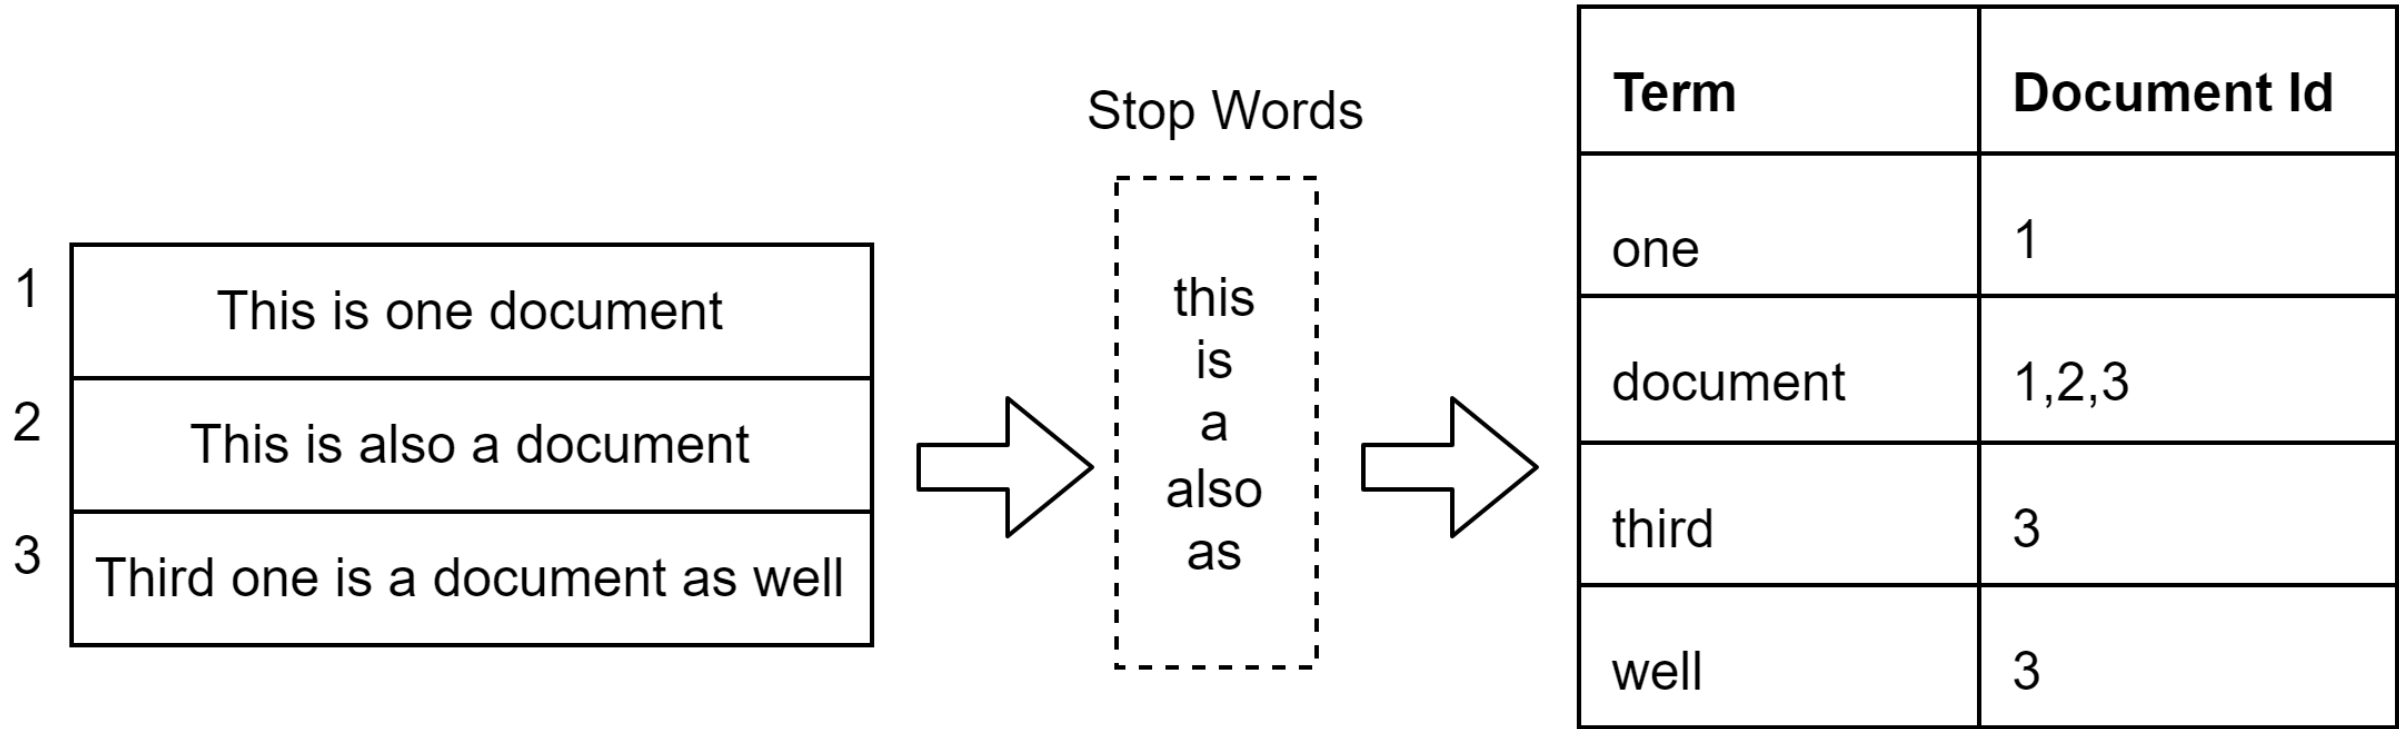


Consuming topics

- Consumers join consumer groups by using the same **group.id**
- The maximum parallelism of a consumer group is the number of partitions for the consumed topic
 - Number of consumers > number of topic partitions will leave consumers idle
- Consume topics through Sink Connectors and the Consume API
- The internal topic `_consumer_offsets` keeps track of each groups current offset
 - Consumers can themselves reset their own offset to any position
- **Kafka Streams** lets us build a streaming application which transforms an input topic into new topics
 - Aggregating, filtering, grouping, joins ++
- **ksqlDB** allows us to perform stream processing tasks using SQL statements

Elastic Stack

- Consists of four tools, **Elasticsearch**, **Logstash**, **Kibana** and **Beats**
- Elasticsearch is a Lucene-based search and analytics engine
- It is distributed and RESTful
- Useful for searching great amounts of data in near-real time
- Documents are indexed and stored in **indices**
 - The generated **inverted indices** tells Elasticsearch which words appear in which document



Ingesting data into Elasticsearch

- REST API
- Logstash can be used to process each incoming message
 - Powerful and flexible tool
 - Higher hardware requirements than Beats
- Beats has been introduced as lightweight data shipper
- Both Logstash and Beats can be used in combination
- In the demo we will be ingesting data using the Confluent Elasticsearch Sink Connector
 - Very easy to setup and fully-managed in Confluent Cloud



Visualizing data

- Kibana is a data visualization dashboard software
- Provides a UI to explore the data in the Elasticsearch indices
- Has a number of different features included
 - Metrics
 - Charts
 - Maps
 - Anomaly detection
 - ++



kibana

Combining Apache Kafka and Elastic Stack

- Elasticsearch can be queried through a REST API
 - Elasticsearch clients are available for a lot of different programming languages
- Make real-time events searchable
- Visualize real-time events using Kibana
- Very easy to consume Kafka topics using the fully-managed Elasticsearch Sink Connector in Confluent Cloud
 - Many more fully-managed and self-managed connectors are available
 - It is also possible to configure Logstash to consume topics, and preprocess messages before they are stored

Demo



Forte
Digital



WWW.FORTEDIGITAL.NO